

Project Initialization and Planning Phase

Date	29 June 2026
Team ID	
Project Name	Personalized Networking Assistant
Maximum Marks	5 Marks

Project Proposal (Proposed Solution) Report:

The proposal report aims to reduce networking preparation friction using AI, boosting attendee confidence and conversation quality at professional events. It tackles the inefficiency of manual, last-minute prep, promising faster preparation, more relevant conversations, and reduced social anxiety. Key features include automated theme extraction, fact-grounded starter generation, and session history with feedback logging.

Project Overview	
Objective	The primary objective is to reduce networking preparation time and social anxiety by automatically generating fact-grounded, personalized conversation starters from event descriptions and attendee bios.
Scope	The project extracts key themes from event descriptions using DistilBERT zero-shot classification, retrieves verified supporting facts via the Wikipedia API, and generates natural-language conversation starters using GPT-2 delivered through a FastAPI backend and Streamlit frontend, with local JSON-based history and feedback logging.
Problem Statement	
Description	Professionals attending networking events, particularly students, introverts, and early-career attendees struggle to prepare meaningful, event-specific conversation openers, relying on manual, last-minute, and often superficial preparation.
Impact	This leads to increased social anxiety, missed relationship-building opportunities, and disengagement from valuable professional events.
Proposed Solution	
Approach	Combine NLP-based theme extraction, real-world fact grounding, and generative text to produce conversation starters that are both personalized and verifiably accurate with facts retrieved before generation so starters are fact-checked by construction, reducing preparation time from potentially hours to under a minute.
Key Features	- DistilBERT-based zero-shot theme extraction from event descriptions - Wikipedia-grounded fact verification integrated directly into starter generation - GPT-2-based natural language starter generation - Session history and thumbs up/down feedback logging - Standalone quick fact-check tool for independent research

Resource Requirements:

Resource	Description	Specification/Allocation
Hardware		
Computing Resources	CPU/GPU Specifications	CPU only (AMD Ryzen 5 7520U). No GPU required, DistilBERT and GPT-2 Small run locally on CPU within acceptable latency for a demo
Memory	RAM Specification	Minimum 4GB, 8GB recommended
Storage	Disk space for data, models, and logs	10GB free disk space — model weights, dependencies, and local JSON data
Software		
Frameworks	Python frameworks	FastAPI, Streamlit
Libraries	Additional Libraries	Hugging Face Transformers, wikipedia-api (Python wrapper), pytest, httpx
Development Environments	IDE	Visual Studio Code / PyCharm
Data		
Data	Source, size, format	No external dataset required, event descriptions and bios are user-entered at runtime; theme classification uses a predefined candidate-label list rather than a trained/fine-tuned dataset